

# INLINING POSTGRES FUNCTIONS NOW AND THEN

Paul A. Jungwirth

12 May 2025

# LINEUP

- Inlining SQL Set-Returning Functions
- Inlining Non-SRF SQL Functions
- Inlining Non-SQL Set-Returning Functions

# SRFS

```
CREATE OR REPLACE FUNCTION visible_sales(user_id INT)
RETURNS SETOF sales
AS $$
    SELECT * FROM sales
    WHERE vendor_id IN (SELECT company_id
                        FROM memberships
                        WHERE user_id = $1);
$$ LANGUAGE sql STABLE;
```

# INLINING SQL SRFS

```
1  =# EXPLAIN ANALYZE SELECT  *
2  FROM    visible_sales_slow(1) AS s
3  WHERE   vendor_id = 5000;
4          QUERY PLAN
5  -----
6  Function Scan on visible_sales_slow s
7    (cost=0.25..12.75 rows=5 width=56)
8    (actual time=57.415..57.670 rows=2 loops=1)
9    Filter: (vendor_id = 5000)
10   Rows Removed by Filter: 51688
11   Planning Time: 0.129 ms
12   Execution Time: 57.925 ms
13  (5 rows)
```

# INLINING SQL SRFS

```
1  =# EXPLAIN ANALYZE SELECT  *
2  FROM    visible_sales_slow(1) AS s
3  WHERE   vendor_id = 5000;
4          QUERY PLAN
5  -----
6   Function Scan on visible_sales_slow s
7     (cost=0.25..12.75 rows=5 width=56)
8     (actual time=57.415..57.670 rows=2 loops=1)
9     Filter: (vendor_id = 5000)
10    Rows Removed by Filter: 51688
11   Planning Time: 0.129 ms
12   Execution Time: 57.925 ms
13  (5 rows)
```

# INLINING SQL SRFS

```
1  =# EXPLAIN ANALYZE SELECT  *
2  FROM    visible_sales_slow(1) AS s
3  WHERE   vendor_id = 5000;
4          QUERY PLAN
5  -----
6  Function Scan on visible_sales_slow s
7    (cost=0.25..12.75 rows=5 width=56)
8    (actual time=57.415..57.670 rows=2 loops=1)
9    Filter: (vendor_id = 5000)
10   Rows Removed by Filter: 51688
11   Planning Time: 0.129 ms
12   Execution Time: 57.925 ms
13  (5 rows)
```

# INLINING SQL SRFS

```
1  =# EXPLAIN ANALYZE SELECT  *
2  FROM    visible_sales_slow(1) AS s
3  WHERE   vendor_id = 5000;
4          QUERY PLAN
5  -----
6  Function Scan on visible_sales_slow s
7    (cost=0.25..12.75 rows=5 width=56)
8    (actual time=57.415..57.670 rows=2 loops=1)
9    Filter: (vendor_id = 5000)
10   Rows Removed by Filter: 51688
11   Planning Time: 0.129 ms
12   Execution Time: 57.925 ms
13  (5 rows)
```

# INLINING SQL SRFS

```
1  =# EXPLAIN ANALYZE SELECT  *
2  FROM    visible_sales_slow(1) AS s
3  WHERE   vendor_id = 5000;
4          QUERY PLAN
5  -----
6  Function Scan on visible_sales_slow s
7    (cost=0.25..12.75 rows=5 width=56)
8    (actual time=57.415..57.670 rows=2 loops=1)
9    Filter: (vendor_id = 5000)
10   Rows Removed by Filter: 51688
11   Planning Time: 0.129 ms
12   Execution Time: 57.925 ms
13  (5 rows)
```



# INLINING SQL SRFS

```
1  Nested Loop
2    (cost=0.84..32.53 rows=7 width=27)
3    (actual time=0.044..0.048 rows=2 loops=1)
4    -> Index Only Scan using idx_memberships_company_user c
5          (cost=0.42..4.44 rows=1 width=4)
6          (actual time=0.025..0.026 rows=1 loops=1)
7          Index Cond: ((company_id = 5000) AND (user_id = 1))
8    -> Index Scan using uq_sales_po_number on sales
9          (cost=0.42..28.02 rows=7 width=27)
10         (actual time=0.013..0.016 rows=2 loops=1)
11         Index Cond: (vendor_id = 5000)
12  Planning Time: 0.258 ms
13  Execution Time: 0.071 ms
```

# INLINING SQL SRFS

```
1  Nested Loop
2    (cost=0.84..32.53 rows=7 width=27)
3    (actual time=0.044..0.048 rows=2 loops=1)
4    -> Index Only Scan using idx_memberships_company_user c
5          (cost=0.42..4.44 rows=1 width=4)
6          (actual time=0.025..0.026 rows=1 loops=1)
7          Index Cond: ((company_id = 5000) AND (user_id = 1))
8    -> Index Scan using uq_sales_po_number on sales
9          (cost=0.42..28.02 rows=7 width=27)
10         (actual time=0.013..0.016 rows=2 loops=1)
11         Index Cond: (vendor_id = 5000)
12  Planning Time: 0.258 ms
13  Execution Time: 0.071 ms
```

# INLINING SQL SRFS

```
1  Nested Loop
2    (cost=0.84..32.53 rows=7 width=27)
3    (actual time=0.044..0.048 rows=2 loops=1)
4    -> Index Only Scan using idx_memberships_company_user c
5          (cost=0.42..4.44 rows=1 width=4)
6          (actual time=0.025..0.026 rows=1 loops=1)
7          Index Cond: ((company_id = 5000) AND (user_id = 1))
8    -> Index Scan using uq_sales_po_number on sales
9          (cost=0.42..28.02 rows=7 width=27)
10         (actual time=0.013..0.016 rows=2 loops=1)
11         Index Cond: (vendor_id = 5000)
12  Planning Time: 0.258 ms
13  Execution Time: 0.071 ms
```

# INLINING SQL SRFS

```
1  Nested Loop
2    (cost=0.84..32.53 rows=7 width=27)
3    (actual time=0.044..0.048 rows=2 loops=1)
4    -> Index Only Scan using idx_memberships_company_user c
5          (cost=0.42..4.44 rows=1 width=4)
6          (actual time=0.025..0.026 rows=1 loops=1)
7          Index Cond: ((company_id = 5000) AND (user_id = 1))
8    -> Index Scan using uq_sales_po_number on sales
9          (cost=0.42..28.02 rows=7 width=27)
10         (actual time=0.013..0.016 rows=2 loops=1)
11         Index Cond: (vendor_id = 5000)
12  Planning Time: 0.258 ms
13  Execution Time: 0.071 ms
```

# INLINING SQL SRFS

```
1  =# EXPLAIN (ANALYZE) SELECT  *
2  FROM    visible_sales(1) AS s
3  WHERE   vendor_id = 2 LIMIT 10;
4          QUERY PLAN
5  -----
6  Limit   (cost=0.42..5.84 rows=10 width=27)
7          (actual time=0.972..1.010 rows=10 loops=1)
8    ->    Nested Loop Semi Join
9          (cost=0.42..2722.73 rows=5020 width=27)
10         (actual time=0.970..1.004 rows=10 loops=1)
11         ->    Seq Scan on sales
12             (cost=0.00..2655.54 rows=5020 width=27)
13             (actual time=0.868..0.896 rows=10 loops=1)
14             Filter: (vendor_id = 2)
15             Rows Removed by Filter: 134
```

# INLINING SQL SRFS

```
1  =# EXPLAIN (ANALYZE) SELECT  *
2  FROM    visible_sales(1) AS s
3  WHERE   vendor_id = 2 LIMIT 10;
4          QUERY PLAN
5  -----
6  Limit    (cost=0.42..5.84 rows=10 width=27)
7           (actual time=0.972..1.010 rows=10 loops=1)
8     ->    Nested Loop Semi Join
9           (cost=0.42..2722.73 rows=5020 width=27)
10          (actual time=0.970..1.004 rows=10 loops=1)
11          ->    Seq Scan on sales
12              (cost=0.00..2655.54 rows=5020 width=27)
13              (actual time=0.868..0.896 rows=10 loops=1)
14              Filter: (vendor_id = 2)
15              Rows Removed by Filter: 134
```

# INLINING SQL SRFS

```
10      (actual time=0.970..1.004 rows=10 loops=1)
11      -> Seq Scan on sales
12            (cost=0.00..2655.54 rows=5020 width=27)
13            (actual time=0.868..0.896 rows=10 loops=1)
14            Filter: (vendor_id = 2)
15            Rows Removed by Filter: 134
16      -> Materialize (cost=0.42..4.44 rows=1 width=4)
17            (actual time=0.010..0.010 rows=1 loops=1)
18      -> Index Only Scan using idx_memberships_co
19            (cost=0.42..4.44 rows=1 width=4)
20            (actual time=0.092..0.092 rows=1 loops=1)
21            Index Cond: ((company_id = 2) AND (use
22 Planning Time: 0.515 ms
23 Execution Time: 1.058 ms
24 (11 rows)
```

# INLINING SQL SRFS

```
10      (actual time=0.970..1.004 rows=10 loops=1)
11      -> Seq Scan on sales
12            (cost=0.00..2655.54 rows=5020 width=27)
13            (actual time=0.868..0.896 rows=10 loops=1)
14            Filter: (vendor_id = 2)
15            Rows Removed by Filter: 134
16      -> Materialize  (cost=0.42..4.44 rows=1 width=4)
17            (actual time=0.010..0.010 rows=1
18      -> Index Only Scan using idx_memberships_co
19            (cost=0.42..4.44 rows=1 width=4)
20            (actual time=0.092..0.092 rows=1 loops
21            Index Cond: ((company_id = 2) AND (use
22  Planning Time: 0.515 ms
23  Execution Time: 1.058 ms
24 (11 rows)
```



# INLINING SQL SRFS

```
10      (actual time=0.970..1.004 rows=10 loops=1)
11      -> Seq Scan on sales
12            (cost=0.00..2655.54 rows=5020 width=27)
13            (actual time=0.868..0.896 rows=10 loops=1)
14            Filter: (vendor_id = 2)
15            Rows Removed by Filter: 134
16      -> Materialize (cost=0.42..4.44 rows=1 width=4)
17            (actual time=0.010..0.010 rows=1 loops=1)
18      -> Index Only Scan using idx_memberships_co
19            (cost=0.42..4.44 rows=1 width=4)
20            (actual time=0.092..0.092 rows=1 loops=1)
21            Index Cond: ((company_id = 2) AND (use
22 Planning Time: 0.515 ms
23 Execution Time: 1.058 ms
24 (11 rows)
```

# temporal\_semijoin

```
SELECT  a.id,  
        UNNEST(multirange(a.valid_at) * j.valid_at) AS valid_at  
FROM    a  
JOIN (     
    SELECT  b.id, range_agg(b.valid_at) AS valid_at  
    FROM    b  
    GROUP BY b.id  
  ) AS j  
ON a.id = j.id AND a.valid_at && j.valid_at;
```

from

# temporal\_semijoin

```
CREATE OR REPLACE FUNCTION temporal_semijoin(  
  left_table text, left_id_col text, left_valid_col text,  
  right_table text, right_id_col text, right_valid_col text  
)  
RETURNS SETOF RECORD AS $$  
DECLARE  
  subquery TEXT := 'j';  
BEGIN  
  IF left_table = 'j' OR right_table = 'j' THEN  
    subquery := 'j1';  
    IF left_table = 'j1' OR right_table = 'j1' THEN  
      subquery := 'j2';  
    END IF;  
  END IF;  
  RETURN QUERY EXECUTE format($j$  
    SELECT %1$I.%2$I, UNNEST(multirange(%1$I.%3$I) * %7$I.%6$I) AS %3$I  
    FROM %1$I  
    JOIN (  
      SELECT %4$I.%5$I, range_agg(%4$I.%6$I) AS %6$I  
      FROM %4$I  
      GROUP BY %4$I.%5$I  
    ) AS %7$I  
    ON %1$I.%2$I = %7$I.%5$I AND %1$I.%3$I && %7$I.%6$I;
```

# INLINING NON-SRFS

```
CREATE OR REPLACE FUNCTION commission_cents(  
  _sale_id INTEGER, _salesperson_id INTEGER  
)  
RETURNS INTEGER  
AS $$  
  SELECT total_price_cents * COALESCE(commission_percent, 0)  
  FROM   sales AS s  
  LEFT JOIN memberships AS m  
  ON     m.company_id = s.vendor_id  
  AND    m.user_id = _salesperson_id  
  WHERE  s.id = _sale_id;  
$$ LANGUAGE sql STABLE;
```

# SUPPORT PROCS

```
1 =# \d pg_proc
```

Table "pg_catalog.pg_proc"					
Column	Type	Collation	Nullable	Default	Storage
oid	oid		not null		
proname	name		not null		
pronamespace	oid		not null		
proowner	oid		not null		
prolang	oid		not null		
procost	real		not null		
prorows	real		not null		
provariadic	oid		not null		
prosupport	regproc		not null		
prokind	"char"		not null		
prosecdef	boolean		not null		

# SUPPORT PROCS

6	proname	name		not null
7	pronamespace	oid		not null
8	proowner	oid		not null
9	prolang	oid		not null
10	procost	real		not null
11	prorows	real		not null
12	provariadic	oid		not null
13	<b>prosupport</b>	<b>regproc</b>		<b>not null</b>
14	prokind	"char"		not null
15	prosecdef	boolean		not null
16	proleakproof	boolean		not null
17	proisstrict	boolean		not null
18	proretset	boolean		not null
19	provolatile	"char"		not null
20	proparallel	"char"		not null

# SUPPORT REQUESTS

- SupportRequestRows
- SupportRequestSelectivity
- SupportRequestCost
- SupportRequestIndexCondition
- SupportRequestWFuncMonotonic
- SupportRequestOptimizeWindowClause
- SupportRequestModifyInPlace
- SupportRequestSimplify

# INLINING NON-SRFS

```
SELECT total_price_cents,  
        commission_cents(id, salesperson_id)  
FROM    sales  
WHERE   salesperson_id = $1  
AND     sold_at BETWEEN start_of_month($2)  
                AND end_of_month($2)
```



# commission\_cents\_support

```
1 CREATE OR REPLACE FUNCTION commission_cents(  
2   _sale_id INTEGER, _salesperson_id INTEGER  
3 )  
4 RETURNS INTEGER  
5 AS $$  
6   SELECT total_price_cents * COALESCE(commission_percent,  
7   FROM sales AS s  
8   LEFT JOIN memberships AS m  
9   ON      m.company_id = s.vendor_id  
10  AND      m.user_id = _salesperson_id  
11  WHERE    s.id = _sale_id;  
12 $$ LANGUAGE sql STABLE  
13 SUPPORT commission_cents_support;
```

# commission\_cents\_support

```
1 CREATE OR REPLACE FUNCTION commission_cents(  
2   _sale_id INTEGER, _salesperson_id INTEGER  
3 )  
4 RETURNS INTEGER  
5 AS $$  
6   SELECT total_price_cents * COALESCE(commission_percent,  
7   FROM sales AS s  
8   LEFT JOIN memberships AS m  
9   ON      m.company_id = s.vendor_id  
10  AND      m.user_id = _salesperson_id  
11  WHERE    s.id = _sale_id;  
12 $$ LANGUAGE sql STABLE  
13 SUPPORT commission_cents_support;
```

# commission\_cents\_support

```
CREATE OR REPLACE FUNCTION commission_cents_support(INTERNAL)  
RETURNS INTERNAL  
AS 'commission_cents', 'commission_cents_support'  
LANGUAGE C;
```

# commission\_cents\_support

```
1 Datum commission_cents_support(PG_FUNCTION_ARGS) {
2     Node *rawreq = (Node *) PG_GETARG_POINTER(0);
3     SupportRequestSimplify *req;
4
5     if (!IsA(req, SupportRequestSimplify)) {
6         PG_RETURN_POINTER(NULL);
7     }
8     req = (SupportRequestSimplify *) rawreq;
9
10    FuncExpr *expr = req->fcall;
11
12    ...
```

# commission\_cents\_support

```
1 Datum commission_cents_support(PG_FUNCTION_ARGS) {  
2     Node *rawreq = (Node *) PG_GETARG_POINTER(0);  
3     SupportRequestSimplify *req;  
4  
5     if (!IsA(req, SupportRequestSimplify)) {  
6         PG_RETURN_POINTER(NULL);  
7     }  
8     req = (SupportRequestSimplify *) rawreq;  
9  
10    FuncExpr *expr = req->fcall;  
11  
12    ...
```

# commission\_cents\_support

```
1 Datum commission_cents_support(PG_FUNCTION_ARGS) {
2     Node *rawreq = (Node *) PG_GETARG_POINTER(0);
3     SupportRequestSimplify *req;
4
5     if (!IsA(req, SupportRequestSimplify)) {
6         PG_RETURN_POINTER(NULL);
7     }
8     req = (SupportRequestSimplify *) rawreq;
9
10    FuncExpr *expr = req->fcall;
11
12    ...
```

# commission\_cents\_support

```
1 Datum commission_cents_support(PG_FUNCTION_ARGS) {
2     Node *rawreq = (Node *) PG_GETARG_POINTER(0);
3     SupportRequestSimplify *req;
4
5     if (!IsA(req, SupportRequestSimplify)) {
6         PG_RETURN_POINTER(NULL);
7     }
8     req = (SupportRequestSimplify *) rawreq;
9
10    FuncExpr *expr = req->fcall;
11
12    ...
```

# commission\_cents\_support

```
typedef struct SupportRequestSimplify
{
    NodeTag      type;

    struct PlannerInfo *root;
    FuncExpr      *fcall;
} SupportRequestSimplify;
```



# commission\_cents\_support

```
1 Node *node = lsecond(expr->args);
2 if (IsA(node, Const)) {
3     Const *c = (Const *) node;
4     if (c->constisnull) {
5         Const *ret = makeConst(
6             INT4OID,          /* type */
7             -1,               /* typmod */
8             0,                /* collid */
9             4,                /* len */
10            Int32GetDatum(0), /* value */
11            false,             /* isnull */
12            true,              /* byval */
13        );
14        PG_RETURN_POINTER(ret);
15    }
```

# commission\_cents\_support

```
1 Node *node = lsecond(expr->args);
2 if (IsA(node, Const)) {
3     Const *c = (Const *) node;
4     if (c->constisnull) {
5         Const *ret = makeConst(
6             INT4OID,          /* type */
7             -1,               /* typmod */
8             0,                /* collid */
9             4,                /* len */
10            Int32GetDatum(0), /* value */
11            false,             /* isnull */
12            true,              /* byval */
13        );
14        PG_RETURN_POINTER(ret);
15    }
```

# commission\_cents\_support

```
1 Node *node = lsecond(expr, args);
2 if (IsA(node, Const)) {
3     Const *c = (Const *) node;
4     if (c->constisnull) {
5         Const *ret = makeConst(
6             INT4OID,          /* type */
7             -1,               /* typmod */
8             0,                /* collid */
9             4,                /* len */
10            Int32GetDatum(0), /* value */
11            false,            /* isnull */
12            true,             /* byval */
13        );
14        PG_RETURN_POINTER(ret);
15    }
16 }
```

# commission\_cents\_support

```
3   Const *c = (Const *) node;
4   if (c->constisnull) {
5       Const *ret = makeConst(
6           INT4OID,          /* type */
7           -1,               /* typmod */
8           0,                /* collid */
9           4,               /* len */
10          Int32GetDatum(0), /* value */
11          false,            /* isnull */
12          true,             /* byval */
13      );
14      PG_RETURN_POINTER(ret);
15  }
16 }
17 PG_RETURN_POINTER(NULL);
```

# THERE'S A PATCH FOR THAT!

TODO: commitfest link? Commit message? Diff snippet?

# SupportRequestInlineSRF

```
1 char *sql = "...";
2 List *parsed = pg_parse_query(sql);
3 List *analyzed = pg_analyze_and_rewrite_with_cb(
4     linitial(parsed),
5     sql,
6     (ParserSetupHook) sql_fn_parser_setup,
7     pinfo,
8     NULL);
9 Query *q = linitial(analyzed);
10 PG_RETURN_POINTER(q);
```

# SupportRequestInlineSRF

```
1 char *sql = "...";
2 List *parsed = pg_parse_query(sql);
3 List *analyzed = pg_analyze_and_rewrite_with_cb(
4     linitial(parsed),
5     sql,
6     (ParserSetupHook) sql_fn_parser_setup,
7     pinfo,
8     NULL);
9 Query *q = linitial(analyzed);
10 PG_RETURN_POINTER(q);
```

# SupportRequestInlineSRF

```
1 char *sql = "...";
2 List *parsed = pg_parse_query(sql);
3 List *analyzed = pg_analyze_and_rewrite_with_cb(
4     linitial(parsed),
5     sql,
6     (ParserSetupHook) sql_fn_parser_setup,
7     pinfo,
8     NULL);
9 Query *q = linitial(analyzed);
10 PG_RETURN_POINTER(q);
```



# SupportRequestInlineSRF

```
1 char *sql = "...";
2 List *parsed = pg_parse_query(sql);
3 List *analyzed = pg_analyze_and_rewrite_with_cb(
4     linitial(parsed),
5     sql,
6     (ParserSetupHook) sql_fn_parser_setup,
7     pinfo,
8     NULL);
9 Query *q = linitial(analyzed);
10 PG_RETURN_POINTER(q);
```

# SupportRequestInlineSRF

```
1 char *sql = "...";
2 List *parsed = pg_parse_query(sql);
3 List *analyzed = pg_analyze_and_rewrite_with_cb(
4     linitial(parsed),
5     sql,
6     (ParserSetupHook) sql_fn_parser_setup,
7     pinfo,
8     NULL);
9 Query *q = linitial(analyzed);
10 PG_RETURN_POINTER(q);
```

**CAVEATS**

# GENERALIZING

```
1 CREATE OR REPLACE FUNCTION temporal_semijoin(  
2   left_table text, left_id_col text, left_valid_col text,  
3   right_table text, right_id_col text, right_valid_col text  
4 )  
5 RETURNS SETOF RECORD AS $$  
6 DECLARE  
7   subquery TEXT := 'j';  
8 BEGIN  
9   IF left_table = 'j' OR right_table = 'j' THEN  
10     subquery := 'j1';  
11     IF left_table = 'j1' OR right_table = 'j1' THEN  
12       subquery := 'j2';  
13     END IF;  
14   END IF;  
15   RETURN QUERY EXECUTE format($$$$  
16     SELECT %1$I.%2$I, UNNEST(multirange(%1$I.%3$I) * %7$I.%6$I) AS %3$I  
17     FROM %1$I  
18     JOIN (  
19       SELECT %4$I.%5$I, range_agg(%4$I.%6$I) AS %6$I  
20       FROM %4$I  
21       GROUP BY %4$I.%5$I  
22     ) AS %7$I  
23     ON %1$I.%2$I = %7$I.%5$I AND %1$I.%3$I && %7$I.%6$I;
```

# GENERALIZING

```
7  subquery TEXT := 'j';
8  BEGIN
9  IF left_table = 'j' OR right_table = 'j' THEN
10     subquery := 'j1';
11     IF left_table = 'j1' OR right_table = 'j1' THEN
12         subquery := 'j2';
13     END IF;
14 END IF;
15 RETURN QUERY EXECUTE format($$$$
16     SELECT  %1$I.%2$I, UNNEST(multirange(%1$I.%3$I) * %7$I.%6$I) AS %3$I
17     FROM    %1$I
18     JOIN (
19         SELECT  %4$I.%5$I, range_agg(%4$I.%6$I) AS %6$I
20         FROM    %4$I
21         GROUP BY %4$I.%5$I
22     ) AS %7$I
23     ON %1$I.%2$I = %7$I.%5$I AND %1$I.%3$I && %7$I.%6$I;
24 $$$,
25 left_table, left_id_col, left_valid_col,
26 right_table, right_id_col, right_valid_col,
27 subquery);
28 END;
29 $$ STABLE LEAKPROOF PARALLEL SAFE SUPPORT temporal_semijoin_support LANGUAGE plpgsql;
```

# THANK YOU!

TODO: github for this talk



# BIBLIOGRAPHY

TODO: github for this talk    TODO: wiki page on sql  
inlining    TODO: commitfest for this patch    TODO:  
temporal\_ops github repo