

# A TEMPORAL BENCHMARK IN BENCHBASE

Paul A. Jungwirth

PGConf.dev 2025

14 May 2025

# TEMPORAL TABLES

```
1 CREATE EXTENSION btree_gist;
2
3 CREATE TABLE employees (
4     id          int GENERATED BY DEFAULT AS IDENTITY,
5     valid_at    daterange,
6     name        text NOT NULL,
7     salary      int NOT NULL,
8     CONSTRAINT employees_pkey
9         PRIMARY KEY (id, valid_at WITHOUT OVERLAPS)
10 );
```

# TEMPORAL TABLES

```
1 CREATE EXTENSION btree_gist;
2
3 CREATE TABLE employees (
4     id          int GENERATED BY DEFAULT AS IDENTITY,
5     valid_at    daterange,
6     name        text NOT NULL,
7     salary      int NOT NULL,
8     CONSTRAINT employees_pkey
9         PRIMARY KEY (id, valid_at WITHOUT OVERLAPS)
10 );
```

# TEMPORAL TABLES

```
1 CREATE EXTENSION btree_gist;
2
3 CREATE TABLE employees (
4     id          int GENERATED BY DEFAULT AS IDENTITY,
5     valid_at    daterange,
6     name        text NOT NULL,
7     salary      int NOT NULL,
8     CONSTRAINT employees_pkey
9         PRIMARY KEY (id, valid_at WITHOUT OVERLAPS)
10 );
```

# TEMPORAL TABLES

```
1 CREATE EXTENSION btree_gist;
2
3 CREATE TABLE employees (
4     id          int GENERATED BY DEFAULT AS IDENTITY,
5     valid_at    daterange,
6     name        text NOT NULL,
7     salary      int NOT NULL,
8     CONSTRAINT employees_pkey
9         PRIMARY KEY (id, valid_at WITHOUT OVERLAPS)
10 );
```

# PAST WORK

**EMPLOYEES**(SSN, LAST\_NAME, FIRST\_NAME, ANNUAL\_SALARY)

**POSITIONS**(PCN, JOB\_TITLE\_CODE1)

**INCUMBENTS**(SSN, PCN, START\_DATE, END\_DATE)

**JOB\_TITLES**(JOB\_TITLE\_CODE, JOB\_TITLE)

Richard Snodgrass,  
University of Arizona employee database

# PAST WORK

# PAST WORK

Our implementation of normalize and alignment. All parameters of the PostgreSQL server, such as maximum memory for sorting, were kept to default values, and no indexes were used.

We use the real world dataset *Incumben* of the University of Arizona with 83,857 entries. Each entry records a job assignment (*pcn*) for an employee (*ssn*) over a specific time interval. The data ranges over 16 years and contains 49,195 different employees. The timestamps are recorded at the granularity of days and range from 1 to 573 days with an average of approximately 180 days. Synthetic datasets used in the evaluation are described below.

## 7.2 Database System Integration

Temporal normalization and temporal alignment fully leverage



# PAST WORK

Our implementation of normalize and alignment fully leverage parameters of the PostgreSQL server, such as maximum memory for sorting, were kept to default values, and no indexes were used.

We use the real world dataset *Incumbent* of the University of Arizona with 83,857 entries. Each entry records a job assignment (*pcn*) for an employee (*ssn*) over a specific time interval. The data ranges over 16 years and contains 49,195 different employees. The timestamps are recorded at the granularity of days and range from 1 to 573 days with an average of approximately 180 days. Synthetic datasets used in the evaluation are described below.

## 7.2 Database System Integration

Temporal normalization and temporal alignment fully leverage

# BENCHBASE NÉE OLTP-BENCH

## OLTP-Bench: An Extensible Testbed for Benchmarking Relational Databases

Djellel Eddine Difallah  
U. of Fribourg, Switzerland  
djelleleddine.difallah@unifr.ch

Carlo Curino  
Microsoft Corporation, USA  
ccurino@microsoft.com

Andrew Pavlo  
Carnegie Mellon University, USA  
pavlo@cs.cmu.edu

Philippe Cudre-Mauroux  
U. of Fribourg, Switzerland  
pcm@unifr.ch

### ABSTRACT

Benchmarking is an essential aspect of any database management system (DBMS) effort. Despite several recent advancements, such

To overcome this problem, it is imperative that application developers use a testing environment that is *stable*, *controlled* and *repeatable* [19]. In the context of DBMSs, this is achieved through the use of a *benchmark* that allows one to measure key performance



# DDL

```
CREATE TABLE employees (  
    id          int GENERATED BY DEFAULT AS IDENTITY NOT NULL,  
    valid_at    daterange NOT NULL,  
    name        text NOT NULL,  
    salary      int NOT NULL,  
    PRIMARY KEY (id, valid_at WITHOUT OVERLAPS)  
);
```

```
CREATE TABLE positions (  
    id          int GENERATED BY DEFAULT AS IDENTITY NOT NULL,  
    valid_at    daterange NOT NULL,  
    name        text NOT NULL,  
    employee_id int NOT NULL,  
    PRIMARY KEY (id, valid_at WITHOUT OVERLAPS),  
    FOREIGN KEY (employee_id, valid_at) REFERENCES empl
```

# InsertPosition

```
1 package com.oltpbenchmark.benchmarks.temporal.procedures;
2
3 public class InsertPosition extends Procedure {
4     public final SQLStmt insertPosition =
5         new SQLStmt(
6             "INSERT INTO positions "
7             + "(employee_id, valid_at, name) "
8             + "VALUES "
9             + "(?, daterange(?, ?), concat(?, ' ', to_char(?,
10             + "RETURNING id"));
11
12     public int run(Connection conn, int employeeId, String du
13         throws SQLException {
14         try (PreparedStatement stmt = this.getPreparedStatement
15             stmt.setInt(1, employeeId);
```

# InsertPosition

```
1 package com.oltpbenchmark.benchmarks.temporal.procedures;
2
3 public class InsertPosition extends Procedure {
4     public final SQLStmt insertPosition =
5         new SQLStmt(
6             "INSERT INTO positions "
7             + "(employee_id, valid_at, name) "
8             + "VALUES "
9             + "(?, daterange(?, ?), concat(?, ' ', to_char(?,
10             + "RETURNING id"));
11
12     public int run(Connection conn, int employeeId, String du
13         throws SQLException {
14         try (PreparedStatement stmt = this.getPreparedStatement
15             stmt.setInt(1, employeeId);
```

# InsertPosition

```
10         + "RETURNING id");
11
12     public int run(Connection conn, int employeeId, String duty,
13         throws SQLException {
14         try (PreparedStatement stmt = this.getPreparedStatement(conn,
15             stmt.setInt(1, employeeId);
16             stmt.setDate(2, s == null ? null : Date.valueOf(s));
17             stmt.setDate(3, e == null ? null : Date.valueOf(e));
18             stmt.setString(4, duty);
19             stmt.setInt(5, rank);
20             stmt.execute();
21             return -1;
22         }
23     }
24 }
```

# PROCEDURES

```
DeleteEmployee.java  
InsertPosition.java  
Noop.java  
SelectAllEmployees.java  
SelectAllEmployeesWithOptionalPositions.java  
SelectEmployeesWithoutPosition.java  
SelectEmployeesWithPosition.java  
SelectOneEmployee.java  
SelectOneEmployeePositions.java  
SelectOneEmployeeWithOptionalPositions.java  
UpdateEmployee.java  
UpdatePosition.java
```



# TemporalWorker

```
private RandomEmployee makeRandomEmployee(boolean gaussianEmpl
    int id = model.gaussianEmployeeId(rng());
    LocalDate s;
    LocalDate e;

    if (TemporalConstants.CHECK_FK_GAUSSIAN_RANGE) {
        s = model.today.plusDays(model.gaussianDays(rng()), (int) M
    } else {
        s = model.today.plusDays(-rng().nextInt(365 * TemporalCons
    }
    // Pick a range from 1 day to 2 years:
    e = s.plusDays(1 + rng().nextInt(365 * 2));

    return new RandomEmployee(id, s, e)
}
```

# FOREIGN KEYS

- EXISTS
- lag
- range\_agg

# EXISTS IMPL

```
SELECT 1
-- There was a PK when the FK started:
WHERE EXISTS
  SELECT 1
  FROM [ONLY] <pktable>
  WHERE pkatt1 = $1 [AND ...]
  AND COALESCE(lower(pkperiodatt), '-Infinity')
    <= COALESCE(lower($n), '-Infinity')
  AND COALESCE(lower($n), '-Infinity')
    < COALESCE(upper(pkperiodatt), 'Infinity')
  FOR KEY SHARE OF x
)
...
```

# EXISTS IMPL

```
-- There was a PK when the FK ended:  
AND EXISTS (  
  SELECT 1  
  FROM    [ONLY] <pktable>  
  WHERE    pkatt1 = $1 [AND ...]  
  AND      COALESCE(lower(pkperiodatt), '-Infinity')  
            < COALESCE(upper($n), 'Infinity')  
  AND      COALESCE(upper($n), 'Infinity')  
            <= COALESCE(upper(pkperiodatt), 'Infinity')  
  FOR KEY SHARE OF x  
)  
...
```

# EXISTS IMPL

```
-- There are no gaps in the PK:  
-- (i.e. there is no PK that ends early,  
-- unless a matching PK record starts right away)  
AND NOT EXISTS (  
  SELECT 1  
  FROM    [ONLY] <pktable> AS pk1  
  WHERE   pkatt1 = $1 [AND ...]  
  AND     COALESCE(lower($n), '-Infinity')  
          < COALESCE(upper(pkperiodatt), 'Infinity')  
  AND     COALESCE(upper(pkperiodatt), 'Infinity')  
          < COALESCE(upper($n), 'Infinity')  
  ...  
)
```

# EXISTS IMPL

```
...
AND      NOT EXISTS (
  SELECT  1
  FROM    [ONLY] <pktable> AS pk2
  WHERE   pk1.pkatt1 = pk2.pkatt1 [AND ...]
          -- but skip pk1.pkperiodatt && pk2.pkperiodatt
  AND     COALESCE(lower(pk2.pkperiodatt), '-Infinity')
          <= COALESCE(upper(pk1.pkperiodatt), 'Infinity')
          COALESCE(upper(pk1.pkperiodatt), 'Infinity')
          <  COALESCE(upper(pk2.pkperiodatt), 'Infinity')
  FOR KEY SHARE OF pk2
)
FOR KEY SHARE OF pk1
)
```

# EXISTS IMPL

```
1 Result (cost=33.28..33.29 rows=1 width=4)
2   One-Time Filter: ((InitPlan 1).col1 AND (InitPlan 2).col1)
3   InitPlan 1
4     -> LockRows (cost=0.28..8.32 rows=1 width=10)
5         -> Index Scan using employees_pkey on employees
6             Index Cond: ((id = 24374) AND (valid_at && 'infinity'))
7             Filter: ((COALESCE(lower(valid_at), '-infinity') < 'infinity'))
8   InitPlan 2
9     -> LockRows (cost=0.28..8.32 rows=1 width=10)
10        -> Index Scan using employees_pkey on employees
11            Index Cond: ((id = 24374) AND (valid_at && 'infinity'))
12            Filter: ((COALESCE(lower(valid_at), '-infinity') < 'infinity'))
13   InitPlan 4
14     -> LockRows (cost=0.28..16.64 rows=1 width=10)
15        -> Index Scan using employees_pkey on employees
```

# EXISTS IMPL

```
9      -> LockRows (cost=0.28..8.32 rows=1 width=10)
10      -> Index Scan using employees_pkey on employees
11          Index Cond: ((id = 24374) AND (valid_at &&
12          Filter: ((COALESCE(lower(valid_at), '-infinity'::date) <=
13  InitPlan 4
14      -> LockRows (cost=0.28..16.64 rows=1 width=10)
15      -> Index Scan using employees_pkey on employees
16          Index Cond: ((id = 24374) AND (valid_at &&
17          Filter: (('2025-05-11'::date < COALESCE(up
18      SubPlan 3
19          -> LockRows (cost=0.28..8.32 rows=1 wi
20          -> Index Scan using employees_pke
21              Index Cond: (id = pk1.id)
22              Filter: ((COALESCE(lower(va
23  (22 rows)
```



# EXISTS IMPL

```
1  Result  (cost=33.28..33.29 rows=1 width=4)
2    One-Time Filter: ((InitPlan 1).col1 AND (InitPlan 2).col1)
3    InitPlan 1
4      -> LockRows  (cost=0.28..8.32 rows=1 width=10)
5          -> Index Scan using employees_pkey on employees
6              Index Cond: ((id = 24374) AND (valid_at && '2000-01-01'))
7              Filter: ((COALESCE(lower(valid_at), '-infinity') < '2000-01-01'))
8    InitPlan 2
9      -> LockRows  (cost=0.28..8.32 rows=1 width=10)
10         -> Index Scan using employees_pkey on employees
11             Index Cond: ((id = 24374) AND (valid_at && '2000-01-01'))
12             Filter: ((COALESCE(lower(valid_at), '-infinity') < '2000-01-01'))
13    InitPlan 4
14      -> LockRows  (cost=0.28..16.64 rows=1 width=10)
15         -> Index Scan using employees_pkey on employees
```

# LAG IMPL

```
1 SELECT 1
2 FROM (
3     SELECT uk.uk_start_value,
4            uk.uk_end_value,
5            NULLIF(LAG(uk.uk_end_value) OVER (ORDER BY uk.uk_start_value), uk.uk_sta
6     FROM (
7         SELECT coalesce(lower(x.pkperiodatt), '-Infinity') AS uk_start_value,
8                coalesce(upper(x.pkperiodatt), 'Infinity') AS uk_end_value
9         FROM pktable AS x
10        WHERE pkatt1 = $1 [AND ...]
11        AND uk.pkperiodatt && $n
12        FOR KEY SHARE OF x
13    ) AS uk
14 ) AS uk
15 WHERE uk.uk_start_value < upper($n)
16 AND uk.uk_end_value >= lower($n)
17 HAVING MIN(uk.uk_start_value) <= lower($n)
18 AND MAX(uk.uk_end_value) >= upper($n)
19 AND array_agg(uk.x) FILTER (WHERE uk.x IS NOT NULL) IS NULL
```

# LAG IMPL

```
1 SELECT 1
2 FROM (
3     SELECT uk.uk_start_value,
4            uk.uk_end_value,
5            NULLIF(LAG(uk.uk_end_value) OVER (ORDER BY uk.uk_start_value), uk.uk_sta
6     FROM (
7         SELECT coalesce(lower(x.pkperiodatt), '-Infinity') AS uk_start_value,
8                coalesce(upper(x.pkperiodatt), 'Infinity') AS uk_end_value
9         FROM pktable AS x
10        WHERE pkatt1 = $1 [AND ...]
11        AND uk.pkperiodatt && $n
12        FOR KEY SHARE OF x
13    ) AS uk
14 ) AS uk
15 WHERE uk.uk_start_value < upper($n)
16 AND uk.uk_end_value >= lower($n)
17 HAVING MIN(uk.uk_start_value) <= lower($n)
18 AND MAX(uk.uk_end_value) >= upper($n)
19 AND array_agg(uk.x) FILTER (WHERE uk.x IS NOT NULL) IS NULL
```

# LAG IMPL

```
1 SELECT 1
2 FROM (
3     SELECT uk.uk_start_value,
4            uk.uk_end_value,
5            NULLIF(LAG(uk.uk_end_value) OVER (ORDER BY uk.uk_start_value), uk.uk_sta
6     FROM (
7         SELECT coalesce(lower(x.pkperiodatt), '-Infinity') AS uk_start_value,
8                coalesce(upper(x.pkperiodatt), 'Infinity') AS uk_end_value
9         FROM pktable AS x
10        WHERE pkatt1 = $1 [AND ...]
11        AND uk.pkperiodatt && $n
12        FOR KEY SHARE OF x
13    ) AS uk
14 ) AS uk
15 WHERE uk.uk_start_value < upper($n)
16 AND uk.uk_end_value >= lower($n)
17 HAVING MIN(uk.uk_start_value) <= lower($n)
18 AND MAX(uk.uk_end_value) >= upper($n)
19 AND array_agg(uk.x) FILTER (WHERE uk.x IS NOT NULL) IS NULL
```

# LAG IMPL

```
1 SELECT 1
2 FROM (
3     SELECT uk.uk_start_value,
4            uk.uk_end_value,
5            NULLIF(LAG(uk.uk_end_value) OVER (ORDER BY uk.uk_start_value), uk.uk_sta
6     FROM (
7         SELECT coalesce(lower(x.pkperiodatt), '-Infinity') AS uk_start_value,
8                coalesce(upper(x.pkperiodatt), 'Infinity') AS uk_end_value
9         FROM   pktable AS x
10        WHERE  pkatt1 = $1 [AND ...]
11        AND    uk.pkperiodatt && $n
12        FOR KEY SHARE OF x
13    ) AS uk
14 ) AS uk
15 WHERE  uk.uk_start_value < upper($n)
16 AND    uk.uk_end_value >= lower($n)
17 HAVING MIN(uk.uk_start_value) <= lower($n)
18 AND    MAX(uk.uk_end_value) >= upper($n)
19 AND    array_agg(uk.x) FILTER (WHERE uk.x IS NOT NULL) IS NULL
```

# LAG IMPL

```
1 SELECT 1
2 FROM (
3     SELECT uk.uk_start_value,
4            uk.uk_end_value,
5            NULLIF(LAG(uk.uk_end_value) OVER (ORDER BY uk.uk_start_value), uk.uk_sta
6     FROM (
7         SELECT coalesce(lower(x.pkperiodatt), '-Infinity') AS uk_start_value,
8                coalesce(upper(x.pkperiodatt), 'Infinity') AS uk_end_value
9         FROM pktable AS x
10        WHERE pkatt1 = $1 [AND ...]
11        AND uk.pkperiodatt && $n
12        FOR KEY SHARE OF x
13    ) AS uk
14 ) AS uk
15 WHERE uk.uk_start_value < upper($n)
16 AND uk.uk_end_value >= lower($n)
17 HAVING MIN(uk.uk_start_value) <= lower($n)
18 AND MAX(uk.uk_end_value) >= upper($n)
19 AND array_agg(uk.x) FILTER (WHERE uk.x IS NOT NULL) IS NULL
```



# range\_agg IMPL

```
1 SELECT 1
2 FROM (
3     SELECT pkperiodatt AS r
4     FROM [ONLY] pktable x
5     WHERE pkatt1 = $1 [AND ...]
6     AND     pkperiodatt && $n
7     FOR KEY SHARE OF x
8 ) x1
9 HAVING $n <@ range_agg(x1.r)
```



# range\_agg IMPL

```
1 SELECT 1
2 FROM (
3     SELECT pkperiodatt AS r
4     FROM [ONLY] pktable x
5     WHERE pkatt1 = $1 [AND ...]
6     AND     pkperiodatt && $n
7     FOR KEY SHARE OF x
8 ) x1
9 HAVING $n <@ range_agg(x1.r)
```

# range\_agg IMPL

```
1 SELECT 1
2 FROM    (
3     SELECT pkperiodatt AS r
4     FROM    [ONLY] pktable x
5     WHERE   pkatt1 = $1 [AND ...]
6     AND     pkperiodatt && $n
7     FOR KEY SHARE OF x
8 ) x1
9 HAVING $n <@ range_agg(x1.r)
```

# range\_agg IMPL

```
1 SELECT 1
2 FROM (
3     SELECT pkperiodatt AS r
4     FROM [ONLY] pktable x
5     WHERE pkatt1 = $1 [AND ...]
6     AND     pkperiodatt && $n
7     FOR KEY SHARE OF x
8 ) x1
9 HAVING $n <@ range_agg(x1.r)
```

# range\_agg IMPL

```
Aggregate (cost=8.32..8.34 rows=1 width=4)
  Filter: ('[2025-05-11,2026-05-26)')::daterange <@ range_agg(
    -> Subquery Scan on x1 (cost=0.28..8.32 rows=1 width=10)
      -> LockRows (cost=0.28..8.31 rows=1 width=16)
        -> Index Scan using employees_pkey on employee
          Index Cond: ((id = 24374) AND (valid_at <= 2025-05-11))
(6 rows)
```

# HYPOTHESIS

- range\_agg fastest
- lag nearly as fast
- EXISTS much slower


# WORKLOAD

```
<transactiontypes>
  <transactiontype>
    <name>InsertPosition</name>
  </transactiontype>
  <transactiontype>
    <name>UpdatePosition</name>
  </transactiontype>
  <transactiontype>
    <name>UpdateEmployee</name>
  </transactiontype>
  <transactiontype>
    <name>DeleteEmployee</name>
  </transactiontype>
</transactiontypes>
```

# PROCEDURES


```
CheckForeignKeyExists.java  
CheckForeignKeyLag.java  
CheckForeignKeyRangeAgg.java
```

# 95% LATENCY





# MEAN LATENCY




# EXISTS IMPL

```
1 Result (actual time=0.020..0.020 rows=0.00 loops=1)
2   One-Time Filter: ((InitPlan 1).col1 AND (InitPlan 2).col1)
3   InitPlan 1
4     -> LockRows (actual time=0.018..0.018 rows=0.00 loops=1)
5         -> Index Scan using employees_pkey on employees
6             Index Cond: ((id = 24374000) AND (valid_at < '2000-01-01'))
7             Filter: ((COALESCE(lower(valid_at), '-infinity') < '2000-01-01'))
8             Index Searches: 1
9   InitPlan 2
10     -> LockRows (never executed)
11         -> Index Scan using employees_pkey on employees
12             Index Cond: ((id = 24374000) AND (valid_at < '2000-01-01'))
13             Filter: ((COALESCE(lower(valid_at), '-infinity') < '2000-01-01'))
14             Index Searches: 0
15   InitPlan 4
```

# EXISTS IMPL

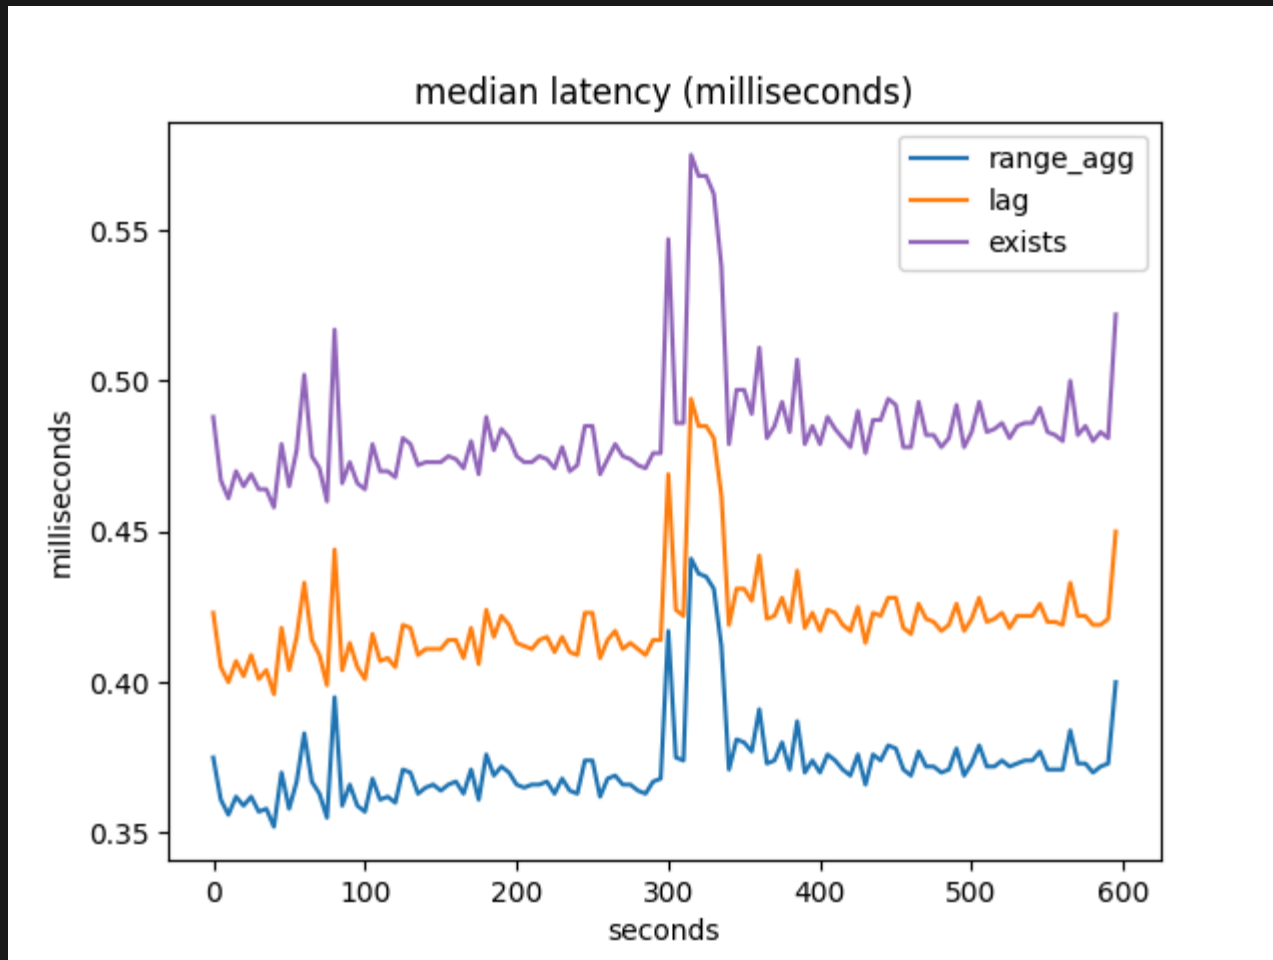
```
0      Index Scan using employees_pkey on employees
9      InitPlan 2
10         -> LockRows (never executed)
11            -> Index Scan using employees_pkey on employees
12               Index Cond: ((id = 24374000) AND (valid_at
13                  Filter: ((COALESCE(lower(valid_at), '-infi
14                  Index Searches: 0
15      InitPlan 4
16         -> LockRows (never executed)
17            -> Index Scan using employees_pkey on employees
18               Index Cond: ((id = 24374000) AND (valid_at
19                  Filter: (('2025-05-11'::date < COALESCE(up
20                  Index Searches: 0
21      SubPlan 3
22         -> LockRows (never executed)
23            -> Index Scan using employees_pkey on employees
```

# 95% LATENCY




100k employees

# MEDIAN LATENCY



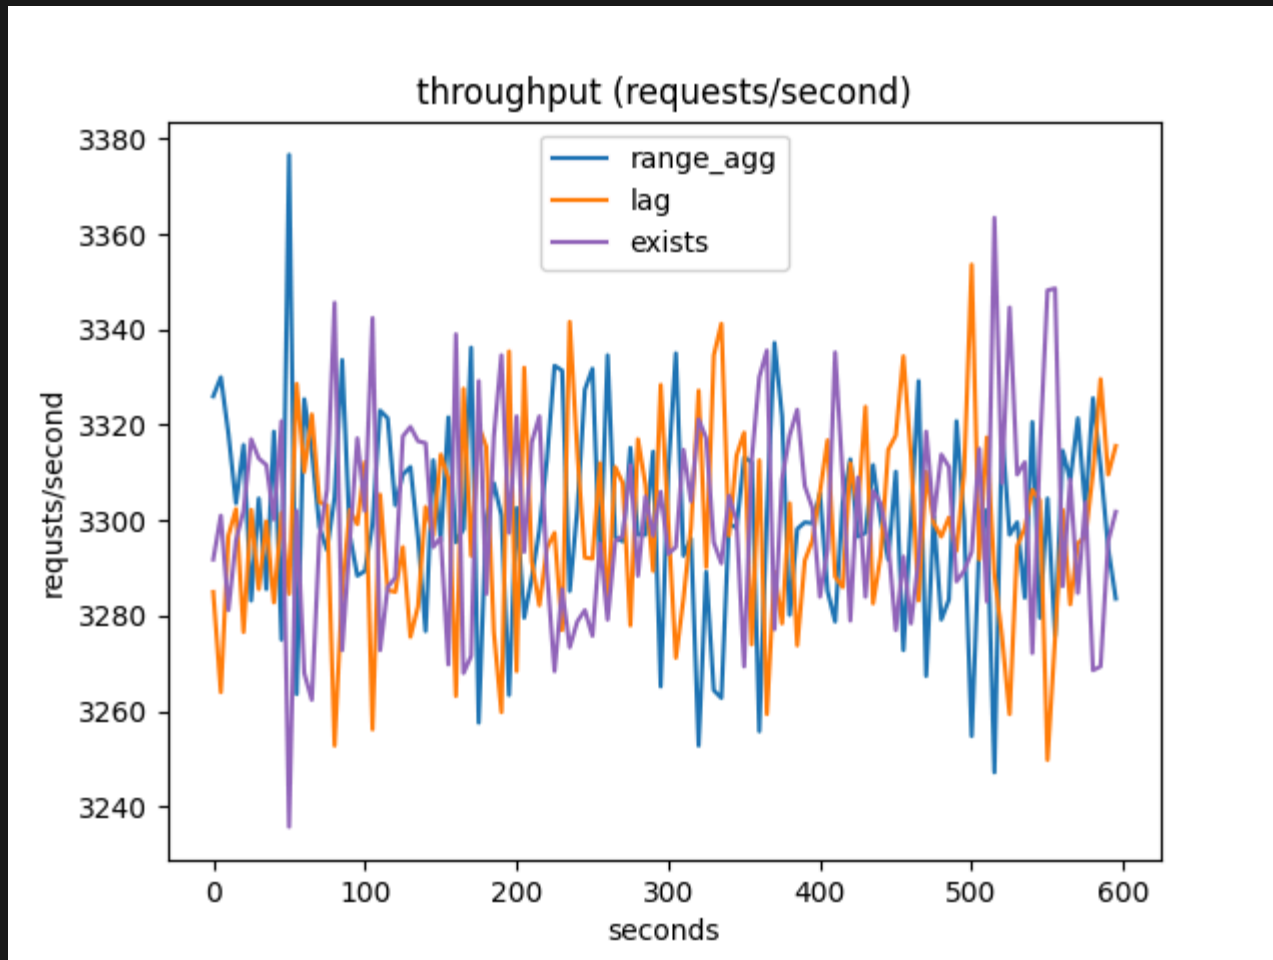
100k employees

# MEAN LATENCY



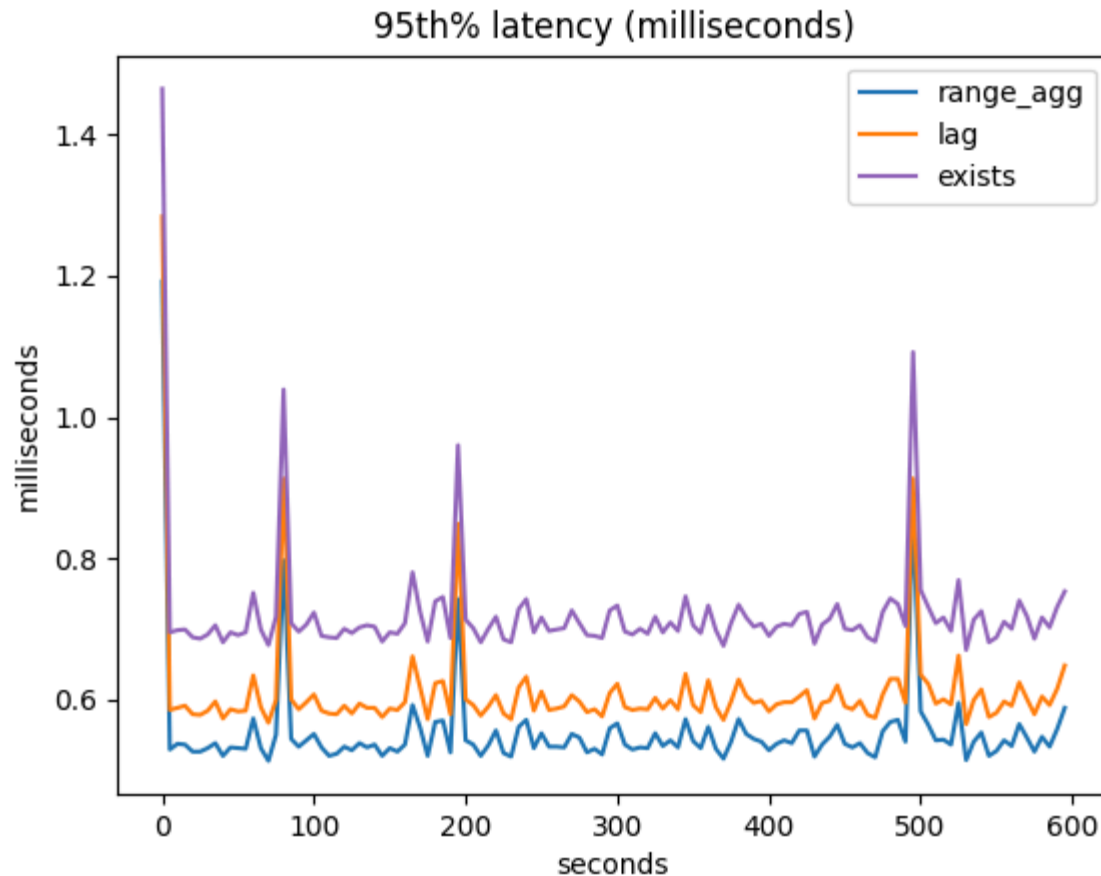
100k employees

# THROUGHPUT



100k employees


# 95% LATENCY



10M employees




# MEDIAN LATENCY



10M employees

# MEAN LATENCY




10M employees

# WITH HISTORY

```
protected void updateEmployees(Connection conn, int lo, int hi) {
    String sql =
        "UPDATE employees FOR PORTION OF valid_at FROM ? TO ? "
        "SET salary = salary * 1.01 " +
        "WHERE id = ?";
    RandomDistribution.Gaussian gaussian =
        new RandomDistribution.Gaussian(this.rng(), 0, config.ge


    try (PreparedStatement employeeUpdate = conn.prepareStatement(sql)) {
        // For each employee:
        for (int i = lo; i <= hi; i++) {
            int raises = gaussian.nextInt();
            LocalDate s = this.model.today;
            LocalDate e;
```

# 95% LATENCY




100k employees, salary history

# MEDIAN LATENCY



100k employees, salary history

# MEAN LATENCY



100k employees, salary history

# FUTURE WORK

- More experiments
- More system stats
- PERIODs
- Compare to other RDBMSes

# REFERENCES

- <https://github.com/pjungwir/introducing-a-temporal-benchmark-in-benchbase>
- <https://dl.acm.org/doi/10.1145/2213836.2213886>
- <https://www.vldb.org/pvldb/vol7/p277-difallah.pdf>
- <https://github.com/cmu-db/benchbase>
- <https://github.com/pjungwir/benchbase/tree/temporal>
- [https://github.com/pjungwir/temporal\\_ops](https://github.com/pjungwir/temporal_ops)
- <https://illuminatedcomputing.com/pages/pdxpug2024-benchbase-and-temporal-foreign-keys/>
- <https://github.com/pjungwir/postgresql/tree/valid-time>
- <https://github.com/pjungwir/benchmarking-temporal-tables>



# THANK YOU

<https://github.com/pjungwir/introducing-a-temporal-benchmark-in-benchbase>